

# MNIST

September 13, 2023

```
[2]: import pandas as pd
import numpy as np
from tqdm.notebook import tqdm
from keras.preprocessing.image import img_to_array, load_img
import tensorflow as tf
import matplotlib.pyplot as plt
%matplotlib inline
import warnings

warnings.filterwarnings('ignore')
```

```
[1]: # !unzip Train_UQcUa52.zip
```

```
[4]: df = pd.read_csv('train.csv')
df.head()
```

```
[4]:  filename  label
0    0.png    4
1    1.png    9
2    2.png    1
3    3.png    7
4    4.png    3
```

```
[5]: !pwd
```

/content

```
[6]: image_path = 'Images/train/'
```

```
[8]: X = np.array([img_to_array(load_img(image_path+df['filename'][i],
→target_size=(28,28,1), grayscale=True))
                for i in tqdm(range(df.shape[0]))
                ]).astype('float32')
```

```
HBox(children=(FloatProgress(value=0.0, max=49000.0), HTML(value='')))
```

```
[9]: y = df['label']
```


```
[10]: print(X.shape, y.shape)
```

```
(49000, 28, 28, 1) (49000,)
```

```
[11]: image_index = 0
      print(y[image_index])
      plt.imshow(X[image_index].reshape(28,28), cmap='Greys')
```

```
4
```


```
[11]: <matplotlib.image.AxesImage at 0x7f813cbf4e48>
```



```
[12]: image_index = 10
      print(y[image_index])
      plt.imshow(X[image_index].reshape(28,28), cmap='Greys')
```

```
2
```


```
[12]: <matplotlib.image.AxesImage at 0x7f813cb8e668>
```



```
[13]: image_index = 100  
      print(y[image_index])  
      plt.imshow(X[image_index].reshape(28,28), cmap='Greys')
```

7

```
[13]: <matplotlib.image.AxesImage at 0x7f813c629c50>
```



```
[14]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
↳random_state=42, stratify=np.array(y))
```

```
[16]: # x_train[0]
```

```
[17]: x_train /= 255
x_test /= 255
```

```
[19]: # x_train[0]
```

```
[20]: input_shape = (28,28,1)
output_class = 10
```

```
[23]: from keras.models import Sequential
from keras.layers import Dense, Conv2D, Dropout, Flatten, MaxPooling2D

# define the model
model = Sequential()
model.add(Conv2D(28, kernel_size=(3,3), input_shape=input_shape))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Flatten())
model.add(Dense(128, activation=tf.nn.relu))
model.add(Dropout(0.3))
model.add(Dense(output_class, activation=tf.nn.softmax))

model.compile(optimizer='adam', loss='sparse_categorical_crossentropy',
↳metrics='accuracy')
```

```
[24]: # train the model
model.fit(x=x_train, y=y_train, batch_size=32, epochs=30,
↳validation_data=(x_test, y_test))
```

Epoch 1/30

1149/1149 [=====] - 10s 3ms/step - loss: 0.4816 -  
accuracy: 0.8475 - val\_loss: 0.1202 - val\_accuracy: 0.9637

Epoch 2/30

1149/1149 [=====] - 4s 3ms/step - loss: 0.1336 -  
accuracy: 0.9605 - val\_loss: 0.0848 - val\_accuracy: 0.9743

Epoch 3/30

1149/1149 [=====] - 4s 3ms/step - loss: 0.0863 -  
accuracy: 0.9732 - val\_loss: 0.0807 - val\_accuracy: 0.9742

Epoch 4/30

1149/1149 [=====] - 4s 3ms/step - loss: 0.0685 -  
accuracy: 0.9783 - val\_loss: 0.0734 - val\_accuracy: 0.9788

Epoch 5/30  
1149/1149 [=====] - 4s 3ms/step - loss: 0.0543 -  
accuracy: 0.9825 - val\_loss: 0.0690 - val\_accuracy: 0.9809

Epoch 6/30  
1149/1149 [=====] - 4s 3ms/step - loss: 0.0461 -  
accuracy: 0.9844 - val\_loss: 0.0684 - val\_accuracy: 0.9808

Epoch 7/30  
1149/1149 [=====] - 4s 3ms/step - loss: 0.0360 -  
accuracy: 0.9873 - val\_loss: 0.0743 - val\_accuracy: 0.9798

Epoch 8/30  
1149/1149 [=====] - 4s 3ms/step - loss: 0.0318 -  
accuracy: 0.9884 - val\_loss: 0.0733 - val\_accuracy: 0.9811

Epoch 9/30  
1149/1149 [=====] - 4s 3ms/step - loss: 0.0319 -  
accuracy: 0.9891 - val\_loss: 0.0658 - val\_accuracy: 0.9838

Epoch 10/30  
1149/1149 [=====] - 4s 3ms/step - loss: 0.0242 -  
accuracy: 0.9919 - val\_loss: 0.0728 - val\_accuracy: 0.9827

Epoch 11/30  
1149/1149 [=====] - 4s 3ms/step - loss: 0.0218 -  
accuracy: 0.9926 - val\_loss: 0.0815 - val\_accuracy: 0.9818

Epoch 12/30  
1149/1149 [=====] - 4s 3ms/step - loss: 0.0286 -  
accuracy: 0.9895 - val\_loss: 0.0766 - val\_accuracy: 0.9829

Epoch 13/30  
1149/1149 [=====] - 4s 3ms/step - loss: 0.0199 -  
accuracy: 0.9928 - val\_loss: 0.0762 - val\_accuracy: 0.9820

Epoch 14/30  
1149/1149 [=====] - 4s 3ms/step - loss: 0.0239 -  
accuracy: 0.9918 - val\_loss: 0.0754 - val\_accuracy: 0.9836

Epoch 15/30  
1149/1149 [=====] - 4s 3ms/step - loss: 0.0160 -  
accuracy: 0.9938 - val\_loss: 0.0865 - val\_accuracy: 0.9820

Epoch 16/30  
1149/1149 [=====] - 4s 3ms/step - loss: 0.0196 -  
accuracy: 0.9935 - val\_loss: 0.0842 - val\_accuracy: 0.9822

Epoch 17/30  
1149/1149 [=====] - 4s 3ms/step - loss: 0.0152 -  
accuracy: 0.9951 - val\_loss: 0.0825 - val\_accuracy: 0.9828

Epoch 18/30  
1149/1149 [=====] - 4s 3ms/step - loss: 0.0155 -  
accuracy: 0.9943 - val\_loss: 0.0889 - val\_accuracy: 0.9817

Epoch 19/30  
1149/1149 [=====] - 4s 3ms/step - loss: 0.0207 -  
accuracy: 0.9930 - val\_loss: 0.0886 - val\_accuracy: 0.9822

Epoch 20/30  
1149/1149 [=====] - 4s 3ms/step - loss: 0.0122 -  
accuracy: 0.9955 - val\_loss: 0.0958 - val\_accuracy: 0.9822

```

Epoch 21/30
1149/1149 [=====] - 4s 3ms/step - loss: 0.0135 -
accuracy: 0.9957 - val_loss: 0.0986 - val_accuracy: 0.9824
Epoch 22/30
1149/1149 [=====] - 4s 3ms/step - loss: 0.0166 -
accuracy: 0.9949 - val_loss: 0.0987 - val_accuracy: 0.9824
Epoch 23/30
1149/1149 [=====] - 4s 3ms/step - loss: 0.0153 -
accuracy: 0.9949 - val_loss: 0.0917 - val_accuracy: 0.9832
Epoch 24/30
1149/1149 [=====] - 4s 3ms/step - loss: 0.0147 -
accuracy: 0.9950 - val_loss: 0.0967 - val_accuracy: 0.9838
Epoch 25/30
1149/1149 [=====] - 4s 3ms/step - loss: 0.0112 -
accuracy: 0.9957 - val_loss: 0.1057 - val_accuracy: 0.9816
Epoch 26/30
1149/1149 [=====] - 4s 3ms/step - loss: 0.0134 -
accuracy: 0.9959 - val_loss: 0.1024 - val_accuracy: 0.9830
Epoch 27/30
1149/1149 [=====] - 4s 3ms/step - loss: 0.0085 -
accuracy: 0.9968 - val_loss: 0.1256 - val_accuracy: 0.9795
Epoch 28/30
1149/1149 [=====] - 4s 3ms/step - loss: 0.0127 -
accuracy: 0.9958 - val_loss: 0.1099 - val_accuracy: 0.9832
Epoch 29/30
1149/1149 [=====] - 4s 3ms/step - loss: 0.0136 -
accuracy: 0.9952 - val_loss: 0.1043 - val_accuracy: 0.9824
Epoch 30/30
1149/1149 [=====] - 4s 3ms/step - loss: 0.0132 -
accuracy: 0.9959 - val_loss: 0.1162 - val_accuracy: 0.9827

```


[24]: <tensorflow.python.keras.callbacks.History at 0x7f80dc057278>

```

[27]: image_index = 10
      # print("Original output:", y_test[image_index])
      plt.imshow(x_test[image_index].reshape(28,28), cmap='Greys')
      pred = model.predict(x_test[image_index].reshape(1,28,28,1))
      print("Predicted output:", pred.argmax())

```

Predicted output: 1



```
[28]: image_index = 100
# print("Original output:", y_test[image_index])
plt.imshow(x_test[image_index].reshape(28,28), cmap='Greys')
pred = model.predict(x_test[image_index].reshape(1,28,28,1))
print("Predicted output:", pred.argmax())
```

Predicted output: 8

